



# Introductory Seminar of PyTorch for Deep Learning

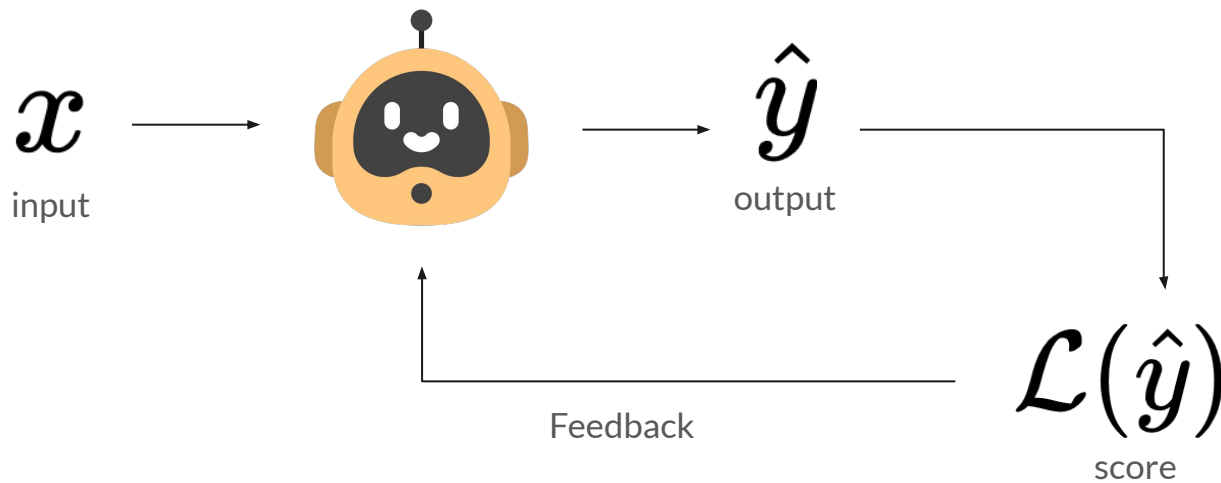
Daniele Angioni, Cagliari Digital Lab 2024 - Day 1



---

# Machine Learning Fundamentals through PyTorch lenses


# Deep Learning Pipeline



# Example



input



output

$$\mathcal{L}(\hat{y})$$


score

Feedback

## Example



input



'giraffe'

output

$$\mathcal{L}(\hat{y})$$

score

Feedback

## Example



input



a giraffe  
standing on  
the grass

output

$$\mathcal{L}(\hat{y})$$

score

Feedback

## Example

a giraffe  
standing on  
the grass

input



output


Feedback

$\mathcal{L}(\hat{y})$   
score

## Example



input



output

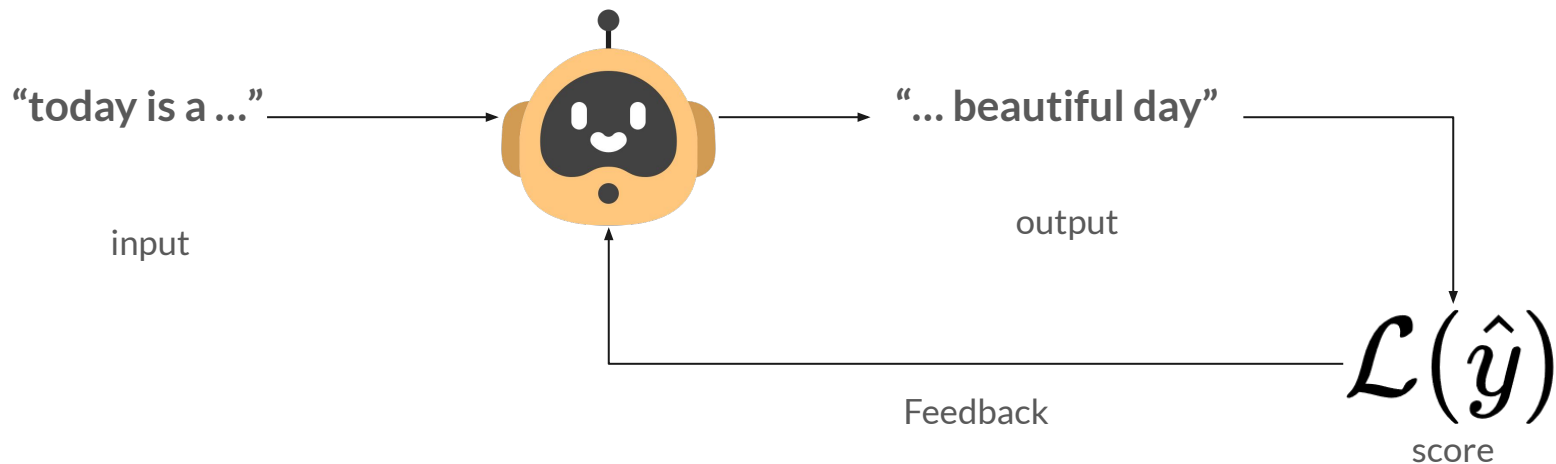
Feedback

$$\mathcal{L}(\hat{y})$$

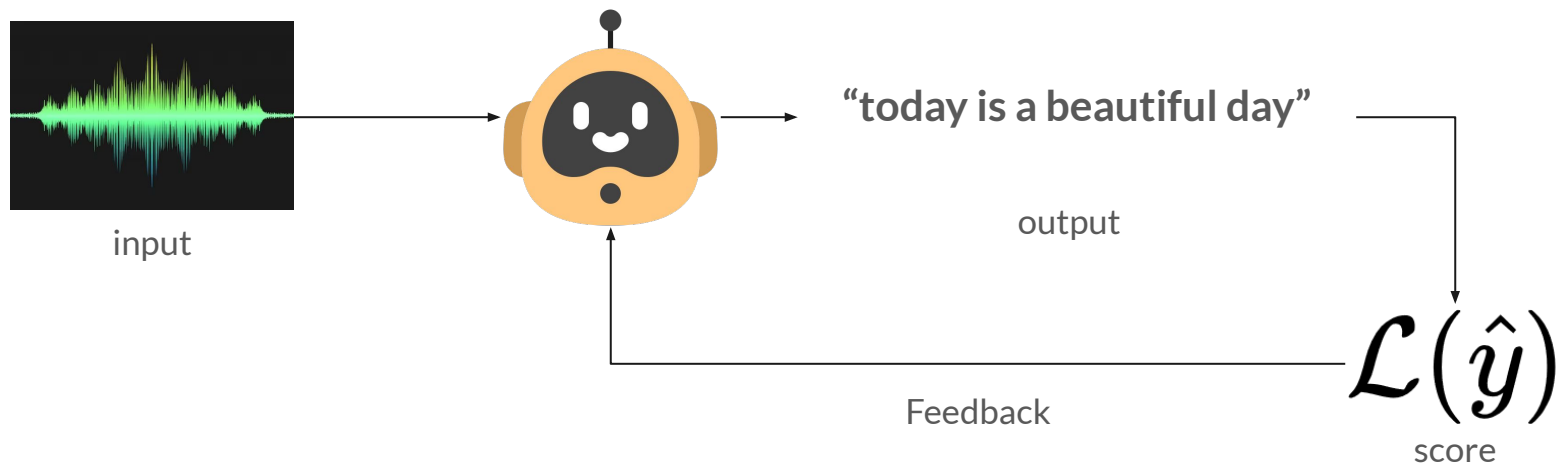
score



## Example



## Example



---


**How can machines deal with such different inputs?**



# Representing data

Deep-learning systems have to be able to map input data to the outputs.

To do so, they usually represent the input data with a specific format.





## What Computer Sees

```

0 2 15 0 0 11 10 0 0 0 0 9 9 0 0 0
0 0 0 4 60 157 236 255 255 177 95 61 32 0 0 29
0 10 16 19 238 255 244 245 243 250 249 255 222 103 10
0 14 170 255 255 244 254 253 253 245 255 249 253 251 124
2 98 255 228 255 251 254 251 141 116 122 215 238 255 49
13 217 243 255 155 133 326 52 2 1 10 13 232 255 255 36
16 229 252 254 49 12 0 0 7 7 0 70 237 252 235 62
6 141 245 255 212 25 11 9 3 0 115 236 255 253 137
0 87 252 250 248 215 60 0 1 121 252 255 248 144 6
0 13 113 255 255 245 255 182 181 248 252 247 206 38 0
0 1 5 117 251 255 241 254 247 255 241 162 17 19
0 0 4 58 251 255 246 254 253 255 120 11 0 1
0 0 0 4 97 255 255 248 252 255 254 255 182 180 4
0 22 206 252 246 251 241 100 1 14 13 255 245 255 194 9
0 111 255 242 255 158 24 0 0 6 39 255 232 230 56
0 218 251 250 137 7 11 0 0 0 2 62 255 250 125 3
0 173 255 255 101 9 20 0 13 3 13 182 251 245 61
0 107 251 241 255 230 98 55 19 118 217 248 253 255 52
0 18 146 255 255 247 255 245 255 249 255 240 255 129 0
0 23 113 255 255 250 248 255 255 245 248 118 14 12
0 0 6 1 0 52 153 233 255 252 147 37 0 0 4 1
0 0 5 0 0 0 0 0 0 14 1 0 0 6 6 0

```

# Audio



# Text

- Vocabulary

"today is a beautiful day"




|           |    |
|-----------|----|
| day       | 1. |
| fire      | 0. |
| beautiful | 1. |
| wood      | 0. |
| is        | 1. |
| arrow     | 0. |
| length    | 0. |
| a         | 1. |
| today     | 1. |
| computer  | 0. |



# Text

- Vocabulary
- Encoding in higher-dimensional space





---

# Tensor Basics



# Deep Learning as Floating Point Numbers

All these data are represented as a collection of floating point numbers structured in a particular way depending on the specific application.

The same can apply to the parameters of a deep learning model, or its output.

These are all collections of floating point numbers!

# Tensors

In the context of deep learning, we use the mathematical concept of **tensor**.  
A tensor is simply the extension of a vector that has an arbitrary number of dimensions.

tensors == multi-dimensional arrays


**Scalar**  
0-dimensional




**Vector**  
1-dimensional



**Matrix**  
2-dimensional



**Tensor**  
N-dimensional





# Tensors in PyTorch

**PyTorch** is not the only library that deals with n-dimensional arrays. **NumPy**, **SciPy**, **Scikit-learn**, **Pandas**, and other deep-learning libraries such as **Tensorflow** also support n-dimensional arrays.



# Tensors in PyTorch

However, in PyTorch, the Tensor class is more powerful than standard numeric libraries.

- GPU support
- Parallel operations on multiple devices or machines
- Keep track of graph of computations that created them

All these features, especially the last one, are of utmost importance when dealing with deep learning!

We will see why in the next chapters...



# Accessing Tensors and their Elements

Tensors are arrays, i.e., **data structures** that store a collection of numbers that are accessible individually using an index, and that can be indexed with multiple indices (at most, one index for each dimension).

# Tensors vs Python lists

Creating a list and accessing one element

```
[1] 1 l = [1, 0, 3]
     2 print(l[0])
```

⇒ 1

Creating a nested list and accessing one element

```
[ ] 1 nested_list = [[0, 1, 2],
                     |   |   |   |   [1, 2, 3]]
     2
     3 print(nested_list[1][2])
```

⇒ 3

## Tensors vs Python lists

```
[4] 1 import torch
     2 t = torch.tensor([[0, 1, 2],
     3 | | | | | [1, 2, 3]])
     4 print(t[0, 1]) # indexing tensors (we will see more indexing tricks later)
```

⇒ tensor(1)

Although on the surface this example doesn't differ much from a list of number objects, under the hood things are completely different.





## How Tensors are Stored in Memory

- Python lists are collections of objects (also of different types) allocated and stored individually in memory
- PyTorch tensors are allocated contiguously in memory blocks containing C numeric types of 32-bit floats.

## Indexing Tensors

```
[4] 1 import torch
    2 t = torch.tensor([[0, 1, 2],
    3 | | | | | [1, 2, 3]])
    4 print(t[0, 1]) # indexing tensors (we will see more indexing tricks later)
```

⇒ tensor(1)

The indexing operation does not create a new tensor by allocating memory and storing the values in it. That would be very inefficient, especially if we had millions of points.

PyTorch indexing directly references the original tensor.

# Fancy Indexing

With tensors, we can use fancy indexing (like [NumPy indexing](#))

```
[ ] 1 x = torch.tensor([0, 1, 2, 3, 4, 5, 6]) # 1-d tensor
    2 element = x[0] # i-th element
    3 first_elements = x[:3] # from start to element 3
    4 last_elements = x[3:] # from element 3 to the end
    5 some_elements = x[3:5] # from element 1 to element 3
```

# Fancy Indexing

Works similarly with 2-d tensors (rows and columns):

```
[ ] 1 x = torch.tensor([[0, 1, 2], [1, 2, 3]]) # 2-d tensor
    2 element = x[0, 0]
    3 row = x[0, :] # works also with x[0] in this case
    4 column = x[:, 0]
    5 some_rows = x[1:, :] # from row 1 to the end, all columns
    6 some_elements = x[1:2, :1] # from row 1 to 2, from column 0 to 1
```

Same applies to n-d tensors as well!



# Tensor element types

Why not using lists or Python numbers?

- The Python interpreter is slow compared to optimized, compiled code.
- PyTorch tensors provide low-level implementations of the data structures and high-level APIs for the operations.
- PyTorch Tensors keep track of the data type in their attribute `dtype`.
- Possible values of `dtype` are: `torch.float32`, `torch.float64`, `torch.int8`, `torch.uint8`, `torch.bool`\*, ...

\*particularly useful for indexing!

## Handling (and changing) tensor dtypes

```
[12] 1 double_precision = torch.tensor([0, 1], dtype=torch.double)
      2 print(double_precision.dtype)
      3 short_tensor = double_precision.short()
      4 print(short_tensor.dtype)
      5 bool_tensor = double_precision.bool()
      6 print(bool_tensor.dtype)
```

```
→ torch.float64
   torch.int16
   torch.bool
```

# Handling (and changing) tensor dtypes

In general you can call the `.type` method and specify the torch data type (a complete list in the [documentation](#))

```
[20] 1 x = torch.tensor([0, 1], dtype=torch.double)
      2 x.type(torch.uint8)

⇒ tensor([0, 1], dtype=torch.uint8)
```

# Basic Tensor operations

## Creation operations and mutations

```
[21] 1 a = torch.ones(3, 2) # 3x2 tensor of only ones
      2 b = torch.zeros(3, 1) # 3x1 tensor of only zeros
      3 c = torch.zeros_like(a) # same shape and type as a
      4 a_t = a.t() # 2x3 tensor (transpose of a)
      5 print(a.shape) # prints the shape (i.e., all the sizes of the dimensions)
```

## Math operations

```
[22] 1 absolute_values = torch.abs(a) # pointwise operations
      2 mean_value = torch.mean(a) # reduction operations
      3 s = a + c # element-wise sum
      4 p = a * c # element-wise product
      5 z = torch.mm(a, c.t()) # matrix multiplication (careful with shapes!)
      6 broadcasting = a + torch.tensor([1, 2]) # torch tries to match shapes
```



## Boolean indexing

Similarly to NumPy, we can use the boolean tensors to select certain elements of another tensor.

```
[11] 1 x = torch.tensor([[ -4, -1, 2], [ 1, -2, 3]]) # 2-d tensor
      2
      3 boolean_mask = (x > 0) # this could be any boolean expression
      4 print(boolean_mask)
      5 print(x[boolean_mask])
```

```
⇒ tensor([[False, False,  True],
          [ True, False,  True]])
   tensor([2, 1, 3])
```

## Tensor storage

```
[ ] 1 a = torch.tensor([1, 2, 3, 4])
    2 b = a[1] # different Tensor, same storage (points to the same location)
    3 c = a.reshape([2, 2]) # same storage, different stride
    4 print(a.storage())
    5 print(c.storage())
    6 print(a.data_ptr() == c.data_ptr()) # same storage
    7 print(c.stride()) # how many storage items to skip for incrementing each dimension
```

**Remember:** the underlying memory is allocated only once, which makes the view operation very lightweight even for large storages.

## Modifying stored values: in-place operations

In-place operations are used to modify directly stored values. The most used one is the `zero_`, that sets to zero all values. They can be recognized by the trailing underscore `_` in their name.

```
[ ] 1 a = torch.ones(3, 2)
     2 a.zero_() # in-place operation, does not create a new tensor
```

The methods that are not in-place, always return a new tensor.



## Moving tensors to the GPU

Moving tensors to the GPU can make computations massively parallel and fast.

Then, all the operations will be performed with GPU operations, while the API remains the same.

PyTorch supports all GPUs that have support for **CUDA** (Compute Unified Device Architecture), a software layer created by Nvidia.

An accelerated version of PyTorch is also available for Apple Silicon, but it is still not very stable.

## Moving tensors to the GPU

Every PyTorch tensor has the attribute `device`, which says where the tensor data is placed in storage. Tensors can be "moved" (rather, copied) to another device by using the method `'to'`.

```
[ ] 1  gpu_tensor = torch.zeros(1, device='cuda') # created on the GPU
    2  cpu_tensor = torch.zeros(1)
    3  to_gpu = cpu_tensor.to(device='cuda') # this creates a copy of the tensor!
    4  to_gpu_another = cpu_tensor.cuda() # shorthand for the previous command
    5  again_to_cpu = to_gpu.cpu() # shorthand for copying the tensor to cpu
```

If your machine has more GPUs, you can also specify which one to use, e.g., `cuda:0`. Note that operations can be performed only between tensors located on the same device.



## Serializing tensors

Until now, we created tensors only in RAM. At some point, we will want to store a tensor in the persistent memory. PyTorch uses pickle to serialize the tensors. Here is how to store a tensor in memory.

```
torch.save(a, 'tensor.pth') # note that the extension is arbitrary
```

And to load back the tensor, a similar API is available.

```
b = torch.load('tensor.pth')
```

---

# PyTorch Autograd Engine



# PyTorch Autograd Engine

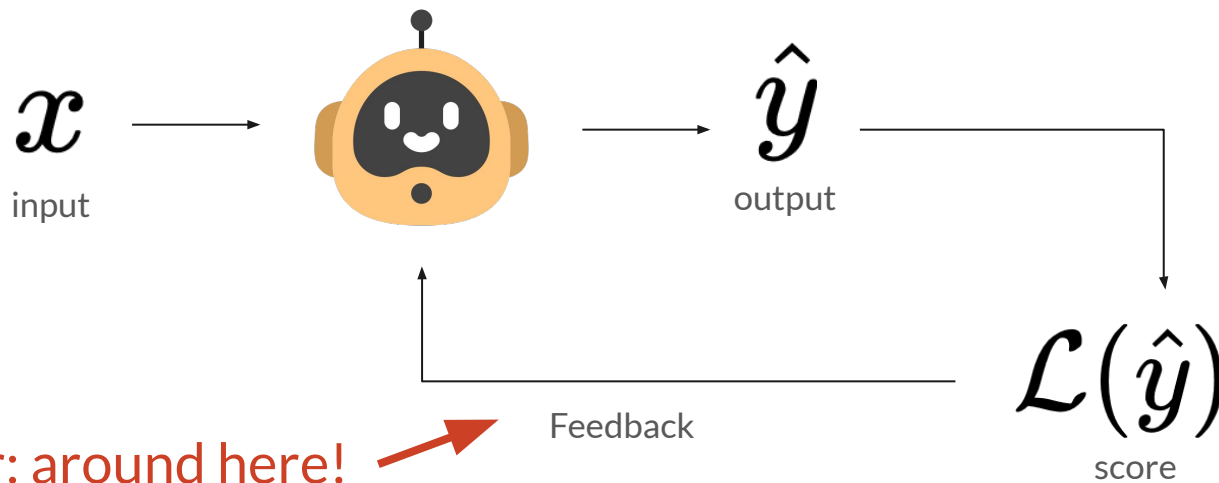
As we will see in the next chapters, in deep learning we need to obtain the **gradients**.

PyTorch computes the gradient of any differentiable function w.r.t. their inputs by using **automatic differentiation** (even for extremely complex functions!)

This PyTorch component is called **autograd**.



## Where DL use gradients?



# The Computational Graph

When operating with tensors, PyTorch automatically build the corresponding **computational graph** by keeping track of every interaction between tensors.

In particular **each node** remembers:

- the **parent tensors** that originated it
- the **operation** performed on the parent tensors





## Getting the gradients (in theory)

In DL we usually want to modify a variable by changing one of the previous ones.

For example, how does modifying the variable **y** influence the variable **x**?

We can find this information in the **partial derivative**  $\frac{\partial y}{\partial x}$

## Getting the gradients (in PyTorch)

In PyTorch, each operation block



$f(x)$

contains the **rule to compute the partial derivative**, and

by calling the `.backward()` method on



$y$

to the gradient field inside



$x$

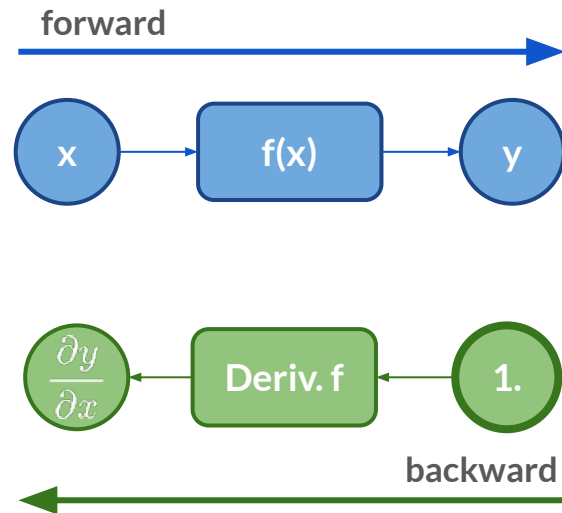
, initialized empty, will be **added the value of the derivative**.

# Forward and backward

In PyTorch, we call **forward pass** the series of operations that start from the inputs to the output nodes of the computational graph.

Instead, we call **backward pass** the series of multiplications of partial derivatives from the root from which we call the `.backward()` method to the inputs (i.e. leaves of the computational graph).

This procedure is nothing else than the implementation of the **backpropagation**, the algorithm that made possible to train neural networks!




## Example

$$y = f(x) = x^2$$

$$\frac{\partial f}{\partial x} = 2x$$

$$\frac{\partial f(x=3)}{\partial x} = ?$$



```
1 x = torch.tensor(3., requires_grad=True)
2 print(x.grad)
3 y = x**2 # forward pass: here we define and use the function
4 print(y)
5 y.backward() # backward pass
6 print(x.grad)
```


```
None
tensor(9., grad_fn=<PowBackward0>)
```

## Example

$$y = f(x) = x^2$$

$$\frac{\partial f}{\partial x} = 2x$$

$$\frac{\partial f(x=3)}{\partial x} = 6$$



```
1 x = torch.tensor(3., requires_grad=True)
2 print(x.grad)
3 y = x**2 # forward pass: here we define and use the function
4 print(y)
5 y.backward() # backward pass
6 print(x.grad)
```

```
None
tensor(9., grad_fn=<PowBackward0>)
tensor(6.)
```

## Accumulated gradients

Remember that every time we call forward and backward, the gradients are not overwritten but accumulated.

This is an important property needed for complex operations, but if disregarded can lead to wrong results.

To fix this issue is sufficient to call the `.zero_()` method on the `grad` attribute

```
[36] 1 y = x**2  
      2 y.backward()  
      3 print(x.grad)
```

↗ tensor(24.)


```
▶ 1 y = x**2  
   2 y.backward()  
   3 print(x.grad)  
   4 x.grad.zero_()
```

↗ tensor(6.)  
 tensor(0.)



## Cascading operations

What if the computational graph starts becoming bigger?



How can we compute the gradient  $\frac{\partial h}{\partial x}$ ?

We know from calculus courses that a cascade of operations can be formally expressed as a **composite function**  $h = g(f(x))$



## Chain Rule

The derivative of a composite function can be computed using the **chain rule**:

$$\frac{\partial h}{\partial x} = \frac{\partial h}{\partial u} \frac{\partial u}{\partial x}$$


If  $u = f(x)$  is the output of  $f(x)$ , then we can compute the derivative of the composed function as the product of the derivative of the "external" function w.r.t.  $u$  and the derivative of the "internal" function w.r.t.  $x$ .

## Example

$$u = f(x) = x^2$$

$$h = g(u) = \log(u)$$

$$\frac{\partial h}{\partial x} = ?$$



```
1 x = torch.tensor(3., requires_grad=True)
2 print(x.grad)
3 y = (x**2).log() # forward (only part of the code that changed)
4 print(y)
5 y.backward() # backward
6 print(x.grad)
```


None  
tensor(2.1972, grad\_fn=<LogBackward0>)

## Example

$$u = f(x) = x^2$$

$$h = g(u) = \log(u)$$

$$\frac{\partial h}{\partial x} = 0.66$$



```
1 x = torch.tensor(3., requires_grad=True)
2 print(x.grad)
3 y = (x**2).log() # forward (only part of the code that changed)
4 print(y)
5 y.backward() # backward
6 print(x.grad)
```

```
None
tensor(2.1972, grad_fn=<LogBackward0>)
tensor(0.6667)
```

## Scaling up...

Thanks to the autograd engine, we can compute gradients from any variable with respect to any other one in an efficient and simple way.

This is of paramount importance when we will deal with large deep learning models!





# Tensor API

Your first source of information should be the [PyTorch documentation](#).

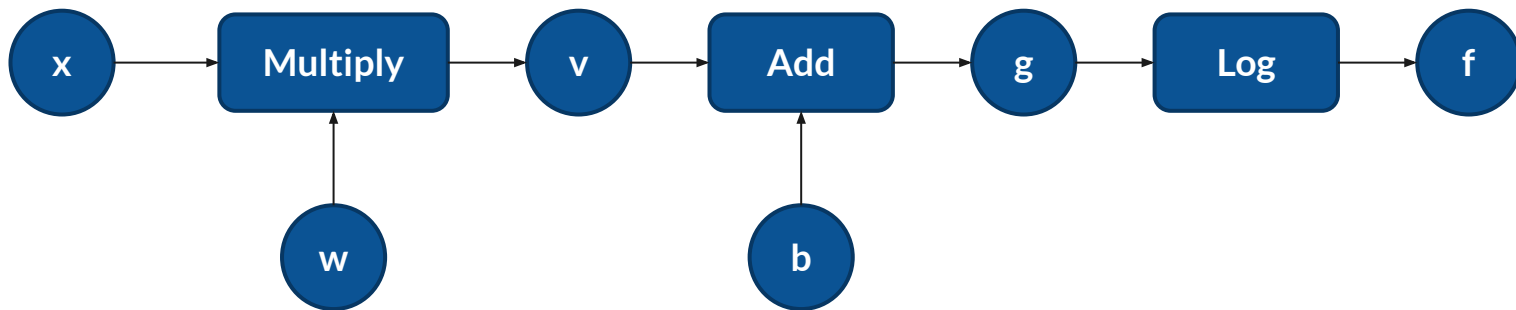
- more complete
- more updated
- (it might also say something different than these slides!)



## Exercise

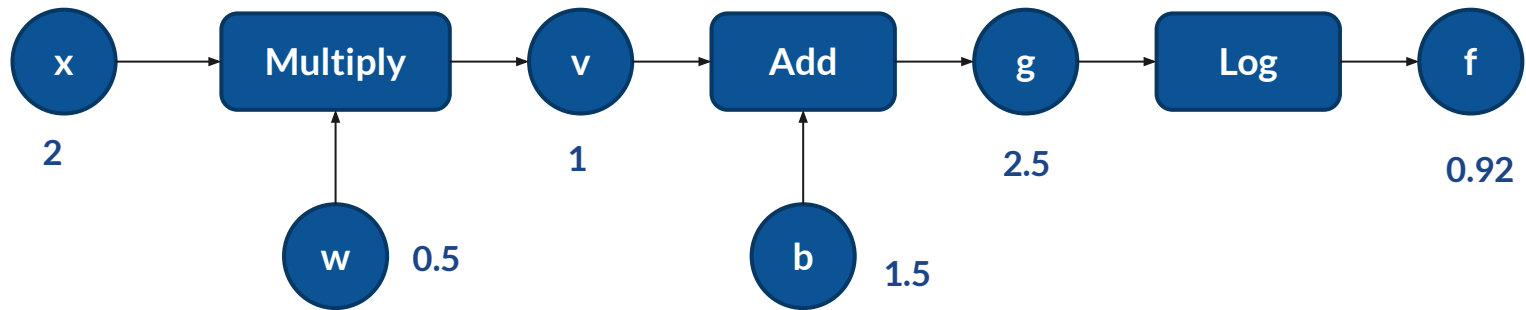
1. Write the **computational graph** for the function  $f = \log(w * x + b)$
2. Compute the **forward** pass setting  $x=2$ ,  $w=0.5$  and  $b=1.5$
3. Compute the **backward** pass **from f to w** (the derivative of  $f$  with respect to  $w$ )

## 1. Computational graph





## 2. Forward pass



### 3. Backward pass

