



# Introductory Seminar of PyTorch for Deep Learning

Daniele Angioni, Cagliari Digital Lab 2024 - Day 5



---

# From PyTorch to Tensorflow



# From PyTorch to Tensforlow

During the seminar we learned the fundamental concepts of Deep Learning to solve different tasks, from computer vision to natural language processing.

In the following we will see first how to apply the fundamentals to solve a different task, and we will also see how to do it with another deep learning framework: Tensorflow

# Tensorflow

Tensorflow is a deep learning framework that, similarly to PyTorch, make use of automatic gradient differentiation.

In Tensorflow we have an high-level API called **Keras**, that makes training and evaluating extremely easy for a user (at least for standard learning problems).



# Keras

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

# Keras pipeline

- Load the dataset

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

# Keras pipeline

- Load the dataset
- Normalization

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

# Keras pipeline

- Load the dataset
- Normalization
- Model definition

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

# Keras pipeline

- Load the dataset
- Normalization
- Model definition
- Setting the optimizer and loss function

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```



# Keras pipeline

- Load the dataset
- Normalization
- Model definition
- Setting the optimizer and loss function
- **Training**

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

# Keras pipeline

- Load the dataset
- Normalization
- Model definition
- Setting the optimizer and loss function
- Training
- Evaluating

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist


(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])


model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```


## Example: Image Denoising




## Solution: using an Autoencoder



## Training an Autoencoder



## Training an Autoencoder



# Loading the MNIST dataset

```
1 import tensorflow as tf
2 import matplotlib.pyplot as plt
3 import numpy as np
4
5 (x_train, _), (x_test, _) = tf.keras.datasets.mnist.load_data()
6 x_train = x_train.astype('float32') / 255.
7 x_test = x_test.astype('float32') / 255.
8
9 # Make sure images have shape (28, 28, 1)
10 x_train = np.expand_dims(x_train, -1)
11 x_test = np.expand_dims(x_test, -1)
12
13 print(x_train.shape)
14 print(x_test.shape)
15 print(type(x_train))
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 [=====] - 2s 0us/step  
(60000, 28, 28, 1)  
(10000, 28, 28, 1)  
<class 'numpy.ndarray'>

## Adding the noise


```
1 noise_factor = 0.3
2 #loc and scale being mean and std
3 x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0,
4                                                         scale=1.0,
5                                                         size=x_train.shape)
6 x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0,
7                                                         scale=1.0,
8                                                         size=x_test.shape)
9
10 #clip pixel under 0 and above 1
11 x_train_noisy = np.clip(x_train_noisy, 0., 1.)
12 x_test_noisy = np.clip(x_test_noisy, 0., 1.)
```



## Visualizing some images

```
1  n_images = 5          #images to be visualized
2  bias = 0              #starting index from the test set for the visualization
3  fig, axs = plt.subplots(nrows=2, ncols=n_images, figsize=(10, 5))
4  for i in range(n_images):
5      axs[0, i].imshow(x_test[bias + i], cmap='gray')
6      axs[0, i].get_xaxis().set_visible(False)
7      axs[0, i].get_yaxis().set_visible(False)
8
9      axs[1, i].imshow(x_test_noisy[bias + i], cmap='gray')
10     axs[1, i].get_xaxis().set_visible(False)
11     axs[1, i].get_yaxis().set_visible(False)
12
13 fig.tight_layout()
14 fig.show()
```

## Visualizing some images



# Model architecture

```
1 from keras.layers import Input, Conv2D, AveragePooling2D, UpSampling2D
2 from keras.models import Model
3 def create_autoencoder():
4     encoder = tf.keras.models.Sequential([
5         Conv2D(32, (5, 5), activation='relu', padding='same'),
6         Conv2D(32, (5, 5), activation='relu', padding='same'),
7         AveragePooling2D((2, 2), padding='same'), #14x14
8         Conv2D(64, (3, 3), activation='relu', padding='same'),
9         Conv2D(64, (3, 3), activation='relu', padding='same'),
10        AveragePooling2D((2, 2), padding='same'), #7x7
11        Conv2D(128, (3, 3), activation='relu', padding='same')
12    ])
13    # latent representation has shape (7, 7, 128)
14    decoder = tf.keras.models.Sequential([
15        Conv2D(128, (3, 3), activation='relu', padding='same'),
16        UpSampling2D((2, 2)), #14x14
17        Conv2D(64, (5, 5), activation='relu', padding='same'),
18        Conv2D(64, (5, 5), activation='relu', padding='same'),
19        UpSampling2D((2, 2)), #28x28
20        Conv2D(32, (5, 5), activation='relu', padding='same'),
21        Conv2D(32, (5, 5), activation='relu', padding='same'),
22        Conv2D(1, (3, 3), activation='relu', padding='same')
23    ])
24    autoencoder = tf.keras.models.Sequential([encoder, decoder])
25    return autoencoder
```

## Training and saving the model weights

```
1  from keras.optimizers import Adam, SGD
2  from keras.losses import MeanSquaredError
3
4  epochs = 10
5  autoencoder = create_autoencoder() #create the architecture
6  optim = SGD(learning_rate=0.001, momentum=0.9)
7  loss = MeanSquaredError()
8  autoencoder.compile(optimizer=optim, loss=loss)
9
10 autoencoder.fit(x_train_noisy, x_train,
11               epochs=10,
12               batch_size=128,
13               shuffle=True,
14               verbose=2,
15               validation_data=(x_test_noisy, x_test))
16 autoencoder.save_weights("autoencoder.h5")
```

## Evaluating

```
1 autoencoder = create_autoencoder() # instantiate the autoencoder model
2 out = autoencoder.predict(x_test[:1]) # build the graph
3 autoencoder.load_weights("autoencoder.h5") # load the pretrained weights
4
5 n_images = 4 #images to be visualized
6 bias = 15 #starting index from the test set for the visualization
7 plt.figure(figsize=(40, 20))
8
9 input_images = x_test_noisy[bias : bias + n_images]
10 target_images = x_test[bias : bias + n_images]
11 output_imgs = autoencoder.predict(x_test_noisy[bias : bias + n_images])
12 output_imgs = np.clip(output_imgs, 0., 1.)
13
14 nrows = 3
15 figdim = 2
16 fig, axs = plt.subplots(nrows=nrows, ncols=n_images, figsize=(figdim*n_images, figdim*nrows))
17
```

# Evaluating

```
18 for i in range(n_images):
19     axs[0, i].imshow(target_images[i], cmap='gray')
20     axs[0, i].set_yticklabels([])
21     axs[0, i].set_xticklabels([])
22
23     axs[1, i].imshow(input_images[i], cmap='gray')
24     axs[1, i].set_yticklabels([])
25     axs[1, i].set_xticklabels([])
26
27     axs[2, i].imshow(output_imgs[i], cmap='gray')
28     axs[2, i].set_yticklabels([])
29     axs[2, i].set_xticklabels([])
30
31 axs[0, 0].set_ylabel('Image')
32 axs[1, 0].set_ylabel('Noisy Image')
33 axs[2, 0].set_ylabel('Reconstructed')
34
35 fig.tight_layout()
36 fig.show()
```

